

# Stroke Prediction Model

KRISHNANDAN SHA

2026-01-09

## About Data Analysis Report

This RMarkdown file contains the report of the data analysis done for the project on building and deploying a stroke prediction model in R. It contains analysis such as data exploration, summary statistics and building the prediction models. The final report was completed on Fri Jan 9 14:20:15 2026.

### Data Description:

According to the World Health Organization (WHO) stroke is the 2nd leading cause of death globally, responsible for approximately 11% of total deaths.

This data set is used to predict whether a patient is likely to get stroke based on the input parameters like gender, age, various diseases, and smoking status. Each row in the data provides relevant information about the patient.

##Task One: Import data and data preprocessing

### Load data and install packages

```
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##     filter, lag
## The following objects are masked from 'package:base':
##     intersect, setdiff, setequal, union
df <- readr::read_csv("healthcare-dataset-stroke-data.csv")

## Rows: 5110 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (6): gender, ever_married, work_type, Residence_type, bmi, smoking_status
## dbl (6): id, age, hypertension, heart_disease, avg_glucose_level, stroke
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
glimpse(df)

## Rows: 5,110
## Columns: 12
```

```
## $ id <dbl> 9046, 51676, 31112, 60182, 1665, 56669, 53882, 10434~  
## $ gender <chr> "Male", "Female", "Male", "Female", "Female", "Male"~  
## $ age <dbl> 67, 61, 80, 49, 79, 81, 74, 69, 59, 78, 81, 61, 54, ~  
## $ hypertension <dbl> 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1~  
## $ heart_disease <dbl> 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0~  
## $ ever_married <chr> "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "No~  
## $ work_type <chr> "Private", "Self-employed", "Private", "Private", "S~  
## $ Residence_type <chr> "Urban", "Rural", "Rural", "Urban", "Rural", "Urban"~  
## $ avg_glucose_level <dbl> 228.69, 202.21, 105.92, 171.23, 174.12, 186.21, 70.0~  
## $ bmi <chr> "36.6", "N/A", "32.5", "34.4", "24", "29", "27.4", "2~  
## $ smoking_status <chr> "formerly smoked", "never smoked", "never smoked", "n~  
## $ stroke <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
```

## Describe and explore the data

```
# Structure of data
str(df)

## #> #> ## #> spc_tbl_ [5,110 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## #> #> $ id : num [1:5110] 9046 51676 31112 60182 1665 ...
## #> #> $ gender : chr [1:5110] "Male" "Female" "Male" "Female" ...
## #> #> $ age : num [1:5110] 67 61 80 49 79 81 74 69 59 78 ...
## #> #> $ hypertension : num [1:5110] 0 0 0 0 1 0 1 0 0 0 ...
## #> #> $ heart_disease : num [1:5110] 1 0 1 0 0 0 1 0 0 0 ...
## #> #> $ ever_married : chr [1:5110] "Yes" "Yes" "Yes" "Yes" ...
## #> #> $ work_type : chr [1:5110] "Private" "Self-employed" "Private" "Private" ...
## #> #> $ Residence_type : chr [1:5110] "Urban" "Rural" "Rural" "Urban" ...
## #> #> $ avg_glucose_level: num [1:5110] 229 202 106 171 174 ...
## #> #> $ bmi : chr [1:5110] "36.6" "N/A" "32.5" "34.4" ...
## #> #> $ smoking_status : chr [1:5110] "formerly smoked" "never smoked" "never smoked" "smokes" ...
## #> #> $ stroke : num [1:5110] 1 1 1 1 1 1 1 1 1 1 ...
## #> - attr(*, "spec")=
## #> .. cols(
## #> ..   id = col_double(),
## #> ..   gender = col_character(),
## #> ..   age = col_double(),
## #> ..   hypertension = col_double(),
## #> ..   heart_disease = col_double(),
## #> ..   ever_married = col_character(),
## #> ..   work_type = col_character(),
## #> ..   Residence_type = col_character(),
## #> ..   avg_glucose_level = col_double(),
## #> ..   bmi = col_character(),
## #> ..   smoking_status = col_character(),
## #> ..   stroke = col_double()
## #> .. )
## #> - attr(*, "problems")=<externalptr>
# Dimensions: rows & columns
dim(df)

## [1] 5110 12

# Column names
colnames(df)
```

```

## [1] "id"                  "gender"          "age"
## [4] "hypertension"        "heart_disease"   "ever_married"
## [7] "work_type"           "Residence_type" "avg_glucose_level"
## [10] "bmi"                 "smoking_status" "stroke"

# Quick tibble-style overview (best for large data)
dplyr::glimpse(df)

## #> #> Rows: 5,110
## #> #> Columns: 12
## #> #> $ id                  <dbl> 9046, 51676, 31112, 60182, 1665, 56669, 53882, 10434~
## #> #> $ gender              <chr> "Male", "Female", "Male", "Female", "Male"~
## #> #> $ age                 <dbl> 67, 61, 80, 49, 79, 81, 74, 69, 59, 78, 81, 61, 54, ~
## #> #> $ hypertension         <dbl> 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1~
## #> #> $ heart_disease       <dbl> 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1~
## #> #> $ ever_married        <chr> "Yes", "Yes", "Yes", "Yes", "Yes", "No~
## #> #> $ work_type            <chr> "Private", "Self-employed", "Private", "Private", "S~
## #> #> $ Residence_type       <chr> "Urban", "Rural", "Rural", "Urban", "Rural", "Urban"~
## #> #> $ avg_glucose_level    <dbl> 228.69, 202.21, 105.92, 171.23, 174.12, 186.21, 70.0~
## #> #> $ bmi                  <chr> "36.6", "N/A", "32.5", "34.4", "24", "29", "27.4", "~
## #> #> $ smoking_status       <chr> "formerly smoked", "never smoked", "never smoked", "~
## #> #> $ stroke               <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
summary(df)

```

```

## #> #> id      gender      age      hypertension
## #> #> Min.   : 67 Length:5110   Min.   : 0.08  Min.   :0.00000
## #> #> 1st Qu.:17741 Class :character 1st Qu.:25.00 1st Qu.:0.00000
## #> #> Median :36932 Mode  :character Median :45.00  Median :0.00000
## #> #> Mean   :36518          Median :43.23  Mean   :0.09746
## #> #> 3rd Qu.:54682          3rd Qu.:61.00  3rd Qu.:0.00000
## #> #> Max.   :72940          Max.   :82.00  Max.   :1.00000
## #> #> heart_disease   ever_married work_type      Residence_type
## #> #> Min.   :0.00000 Length:5110     Length:5110     Length:5110
## #> #> 1st Qu.:0.00000 Class :character  Class :character  Class :character
## #> #> Median :0.00000 Mode  :character  Mode  :character  Mode  :character
## #> #> Mean   :0.05401          Mode  :character  Mean   :0.04873
## #> #> 3rd Qu.:0.00000          Mode  :character  3rd Qu.:0.00000
## #> #> Max.   :1.00000          Max.   :1.00000  Max.   :1.00000
## #> #> avg_glucose_level bmi      smoking_status stroke
## #> #> Min.   : 55.12 Length:5110   Length:5110   Min.   :0.00000
## #> #> 1st Qu.: 77.25 Class :character Class :character 1st Qu.:0.00000
## #> #> Median : 91.89 Mode  :character Mode  :character  Median :0.00000
## #> #> Mean   :106.15          Mode  :character          Mean   :0.04873
## #> #> 3rd Qu.:114.09          Mode  :character          3rd Qu.:0.00000
## #> #> Max.   :271.74          Max.   :271.74   Max.   :1.00000

```

```
#Separate numerical & categorical summaries
```

```

df %>%
  select(where(is.numeric)) %>%
  summary()

## #> #> id      age      hypertension      heart_disease
## #> #> Min.   : 67  Min.   : 0.08  Min.   :0.00000  Min.   :0.00000
## #> #> 1st Qu.:17741 1st Qu.:25.00  1st Qu.:0.00000  1st Qu.:0.00000

```

```

## Median :36932   Median :45.00   Median :0.00000   Median :0.00000
## Mean   :36518   Mean   :43.23   Mean   :0.09746   Mean   :0.05401
## 3rd Qu.:54682   3rd Qu.:61.00   3rd Qu.:0.00000   3rd Qu.:0.00000
## Max.   :72940   Max.   :82.00   Max.   :1.00000   Max.   :1.00000
## avg_glucose_level      stroke
## Min.   : 55.12   Min.   :0.00000
## 1st Qu.: 77.25   1st Qu.:0.00000
## Median : 91.89   Median :0.00000
## Mean   :106.15   Mean   :0.04873
## 3rd Qu.:114.09   3rd Qu.:0.00000
## Max.   :271.74   Max.   :1.00000

#Categorical variables frequency

df %>%
  select(where(is.character)) %>%
  lapply(table)

## $gender
##
## Female   Male   Other
## 2994    2115     1
##
## $ever_married
##
## No   Yes
## 1757  3353
##
## $work_type
##
##       children      Govt_job  Never_worked      Private Self-employed
##             687          657            22           2925            819
##
## $Residence_type
##
## Rural Urban
## 2514  2596
##
## $bmi
##
## 10.3 11.3 11.5   12 12.3 12.8   13 13.2 13.3 13.4 13.5 13.7 13.8 13.9   14 14.1
##   1   1   1   1   1   1   1   1   1   1   1   1   2   2   1   1   5
## 14.2 14.3 14.4 14.5 14.6 14.8 14.9   15 15.1 15.2 15.3 15.4 15.5 15.6 15.7 15.8
##   4   3   2   2   4   4   1   2   8   4   4   3   5   3   3   5
## 15.9 16 16.1 16.2 16.3 16.4 16.5 16.6 16.7 16.8 16.9   17 17.1 17.2 17.3 17.4
##   5   8   8   10  10   11   4   8   11   7   7   10   12   11   9   13
## 17.5 17.6 17.7 17.8 17.9   18 18.1 18.2 18.3 18.4 18.5 18.6 18.7 18.8 18.9   19
##   7   16  13   7   7   16  12   9   17   10   12   19   13   15   8   7
## 19.1 19.2 19.3 19.4 19.5 19.6 19.7 19.8 19.9   20 20.1 20.2 20.3 20.4 20.5 20.6
##   11  13   9   14  21   7   8   17   9   17   25   16   17   23   18   15
## 20.7 20.8 20.9   21 21.1 21.2 21.3 21.4 21.5 21.6 21.7 21.8 21.9   22 22.1 22.2
##   12  17  13   17  16   16   21   22   27   16   14   18   14   15   22   30
## 22.3 22.4 22.5 22.6 22.7 22.8 22.9   23 23.1 23.2 23.3 23.4 23.5 23.6 23.7 23.8
##   18  22  13   18  25   25   16   27   21   20   19   36   31   24   16   22
## 23.9 24 24.1 24.2 24.3 24.4 24.5 24.6 24.7 24.8 24.9   25 25.1 25.2 25.3 25.4

```

```

##   24   28   28   29   26   23   26   22   22   31   27   27   34   18   28   26
## 25.5 25.6 25.7 25.8 25.9   26 26.1 26.2 26.3 26.4 26.5 26.6 26.7 26.8 26.9   27
##   33   21   15   24   24   25   37   27   23   34   30   29   37   21   34   35
## 27.1 27.2 27.3 27.4 27.5 27.6 27.7 27.8 27.9   28 28.1 28.2 28.3 28.4 28.5 28.6
##   28   24   36   22   29   37   37   23   28   28   29   25   30   38   27   27
## 28.7 28.8 28.9   29 29.1 29.2 29.3 29.4 29.5 29.6 29.7 29.8 29.9   30 30.1 30.2
##   41   26   31   26   29   26   22   30   26   26   27   23   26   27   26   20
## 30.3 30.4 30.5 30.6 30.7 30.8 30.9   31 31.1 31.2 31.3 31.4 31.5 31.6 31.7 31.8
##   30   17   24   18   23   21   27   22   26   19   21   30   27   21   14   24
## 31.9 32 32.1 32.2 32.3 32.4 32.5 32.6 32.7 32.8 32.9   33 33.1 33.2 33.3 33.4
##   22   21   24   20   28   19   21   19   21   25   13   15   25   17   15   16
## 33.5 33.6 33.7 33.8 33.9   34 34.1 34.2 34.3 34.4 34.5 34.6 34.7 34.8 34.9   35
##   23   11   19   13   13   17   15   17   18   18   21   11   20   15   11   12
## 35.1 35.2 35.3 35.4 35.5 35.6 35.7 35.8 35.9   36 36.1 36.2 36.3 36.4 36.5 36.6
##   10   16   12   9   13   15   13   24   18   11   7   12   13   10   4   14
## 36.7 36.8 36.9   37 37.1 37.2 37.3 37.4 37.5 37.6 37.7 37.8 37.9   38 38.1 38.2
##   15   8   13   10   7   9   13   11   9   9   7   10   11   13   10   9
## 38.3 38.4 38.5 38.6 38.7 38.8 38.9   39 39.1 39.2 39.3 39.4 39.5 39.6 39.7 39.8
##   2   6   7   9   11   10   8   7   8   10   6   10   8   9   8   5
## 39.9 40 40.1 40.2 40.3 40.4 40.5 40.6 40.7 40.8 40.9   41 41.1 41.2 41.3 41.4
##   5   6   10   10   8   9   7   1   1   7   6   3   7   8   6   3
## 41.5 41.6 41.7 41.8 41.9   42 42.1 42.2 42.3 42.4 42.5 42.6 42.7 42.8 42.9   43
##   8   5   7   11   5   3   3   8   5   5   2   4   4   3   3   8
## 43.1 43.2 43.3 43.4 43.6 43.7 43.8 43.9   44 44.1 44.2 44.3 44.4 44.5 44.6 44.7
##   4   4   5   6   4   7   9   8   4   1   4   3   1   4   2   6
## 44.8 44.9   45 45.1 45.2 45.3 45.4 45.5 45.7 45.8 45.9   46 46.1 46.2 46.3 46.4
##   4   2   5   2   3   4   4   4   2   1   2   4   2   2   1   1
## 46.5 46.6 46.8 46.9 47.1 47.3 47.4 47.5 47.6 47.8 47.9   48 48.1 48.2 48.3 48.4
##   2   1   1   2   1   2   1   3   3   2   1   1   1   1   2   1
## 48.5 48.7 48.8 48.9 49.2 49.3 49.4 49.5 49.8 49.9 50.1 50.2 50.3 50.4 50.5 50.6
##   2   1   2   3   1   3   1   2   3   1   2   4   2   1   1   2
## 50.8 50.9   51 51.5 51.7 51.8 51.9 52.3 52.5 52.7 52.8 52.9 53.4 53.5 53.8 53.9
##   1   1   1   1   1   2   1   1   2   3   1   2   1   2   1   1
## 54 54.1 54.2 54.3 54.6 54.7 54.8   55 55.1 55.2 55.7 55.9   56 56.1 56.6 57.2
##   1   1   1   1   2   3   1   2   1   1   4   2   1   1   2   2
## 57.3 57.5 57.7 57.9 58.1 59.7 60.2 60.9 61.2 61.6 63.3 64.4 64.8 66.8 71.9   78
##   1   1   1   1   1   1   1   2   1   1   1   1   1   1   1   1
## 92 97.6  N/A
##   1   1   201
##
## $smoking_status
##
## formerly smoked    never smoked      smokes      Unknown
##           885          1892          789          1544

#Missing values check (critical for EDA)
colSums(is.na(df))

```

```

##          id      gender       age hypertension
##          0          0          0            0
## heart_disease ever_married work_type Residence_type
##          0          0          0            0
## avg_glucose_level      bmi smoking_status      stroke
##          0          0          0            0

```

```

#Target variable distribution (business-critical)

## Stroke count
table(df$stroke)

##
##      0      1
## 4861  249

## Stroke percentage
prop.table(table(df$stroke)) * 100

##
##          0           1
## 95.127202  4.872798

#Unique values per column

sapply(df, function(x) length(unique(x)))

##             id            gender            age        hypertension
##             5110                  3                 104                      2
##   heart_disease    ever_married       work_type Residence_type
##             2                     2                   5                      2
## avg_glucose_level            bmi   smoking_status        stroke
##             3979                  419                  4                      2

#Quick data quality checklist

## Check duplicates
sum(duplicated(df$id))

## [1] 0

## Check impossible ages
summary(df$age)

##   Min. 1st Qu. Median   Mean 3rd Qu.   Max.
## 0.08  25.00  45.00 43.23  61.00  82.00

#CLEAN & FIX THE DATA
df <- df %>%
  mutate(
    hypertension = factor(hypertension, levels = c(0,1), labels = c("No","Yes")),
    heart_disease = factor(heart_disease, levels = c(0,1), labels = c("No","Yes")),
    stroke       = factor(stroke, levels = c(0,1), labels = c("No","Yes"))
  )

#Fix BMI column
df <- df %>%
  mutate(
    bmi = as.numeric(ifelse(bmi == "N/A", NA, bmi))
  )

```

```

#Handle missing values
colSums(is.na(df))

##          id      gender       age hypertension
##        0         0         0            0
## heart_disease ever_married work_type Residence_type
##        0         0         0            0
## avg_glucose_level      bmi   smoking_status stroke
##        0        201            0            0

#Logical imputation for BMI
median_bmi <- median(df$bmi, na.rm = TRUE)

df <- df %>%
  mutate(
    bmi = ifelse(is.na(bmi), median_bmi, bmi)
  )

#Standardize categorical values (data hygiene)
df <- df %>%
  mutate(
    gender = trimws(gender),
    smoking_status = trimws(smoking_status),
    work_type = trimws(work_type),
    Residence_type = trimws(Residence_type)
  )

#Convert categorical columns to factors (model-ready)
df <- df %>%
  mutate(
    gender      = factor(gender),
    ever_married = factor(ever_married),
    work_type    = factor(work_type),
    Residence_type = factor(Residence_type),
    smoking_status = factor(smoking_status)
  )

#Check duplicates (defensive analytics)
sum(duplicated(df$id))

## [1] 0

#Validation
str(df)

## #tibble [5,110 x 12] (S3: tbl_df/tbl/data.frame)
## $ id           : num [1:5110] 9046 51676 31112 60182 1665 ...
## $ gender       : Factor w/ 3 levels "Female","Male",...: 2 1 2 1 1 2 2 1 1 1 ...
## $ age          : num [1:5110] 67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 1 2 1 1 1 ...
## $ heart_disease: Factor w/ 2 levels "No","Yes": 2 1 2 1 1 1 2 1 1 1 ...
## $ ever_married : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 1 2 2 ...
## $ work_type    : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4 5 4 4 4 4 4 ...
## $ Residence_type: Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2 1 2 ...
## $ avg_glucose_level: num [1:5110] 229 202 106 171 174 ...
## $ bmi          : num [1:5110] 36.6 28.1 32.5 34.4 24 29 27.4 22.8 28.1 24.2 ...

```

```

## $ smoking_status : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1 2 2 4 4 ...
## $ stroke         : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 ...
summary(df)

##      id          gender        age   hypertension heart_disease
##  Min.   : 67   Female:2994   Min.   : 0.08   No :4612    No :4834
##  1st Qu.:17741  Male  :2115   1st Qu.:25.00   Yes: 498    Yes: 276
##  Median :36932  Other :  1   Median :45.00
##  Mean   :36518                   Mean   :43.23
##  3rd Qu.:54682                   3rd Qu.:61.00
##  Max.   :72940                   Max.   :82.00
##      ever_married    work_type Residence_type avg_glucose_level
##  No :1757     children   : 687   Rural:2514      Min.   : 55.12
##  Yes:3353    Govt_job   : 657   Urban:2596     1st Qu.: 77.25
##           Never_worked : 22
##           Private     :2925
##           Self-employed: 819
##           Max.       :271.74
##      bmi          smoking_status stroke
##  Min.   :10.30  formerly smoked: 885   No :4861
##  1st Qu.:23.80  never smoked   :1892   Yes: 249
##  Median :28.10  smokes       : 789
##  Mean   :28.86  Unknown      :1544
##  3rd Qu.:32.80
##  Max.   :97.60

library(dplyr)
library(ggplot2)

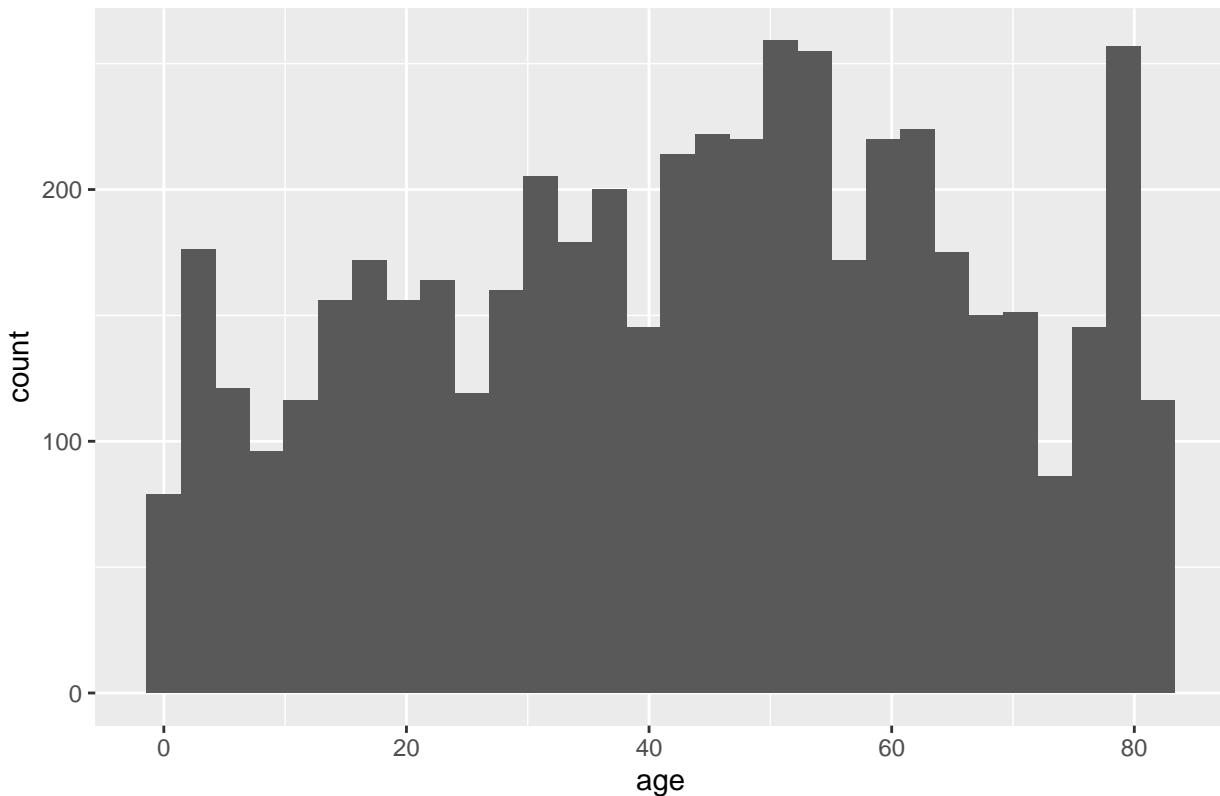
#Age distribution
summary(df$age)

##      Min. 1st Qu. Median   Mean 3rd Qu.   Max.
##  0.08   25.00  45.00  43.23  61.00  82.00

ggplot(df, aes(x = age)) +
  geom_histogram(bins = 30) +
  labs(title = "Age Distribution")

```

## Age Distribution

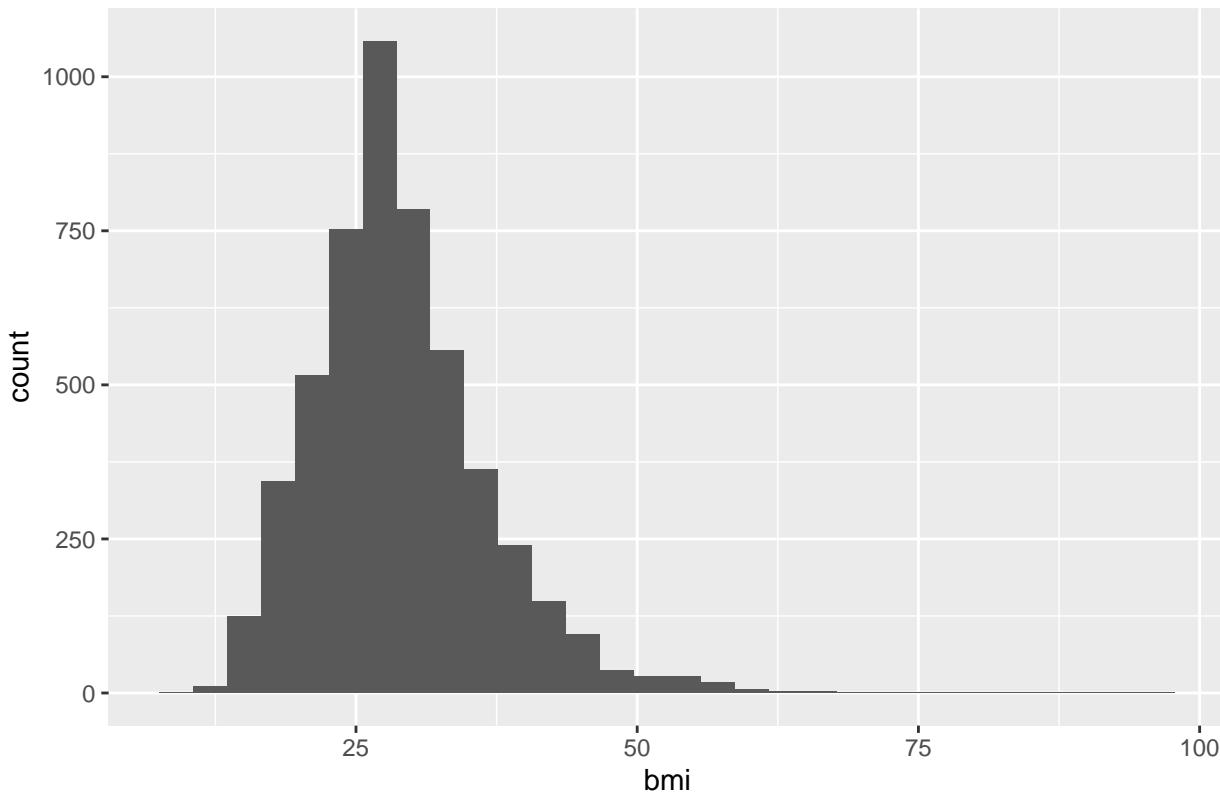


```
#BMI distribution
summary(df$bmi)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    10.30   23.80   28.10   28.86   32.80   97.60

ggplot(df, aes(x = bmi)) +
  geom_histogram(bins = 30) +
  labs(title = "BMI Distribution")
```

## BMI Distribution



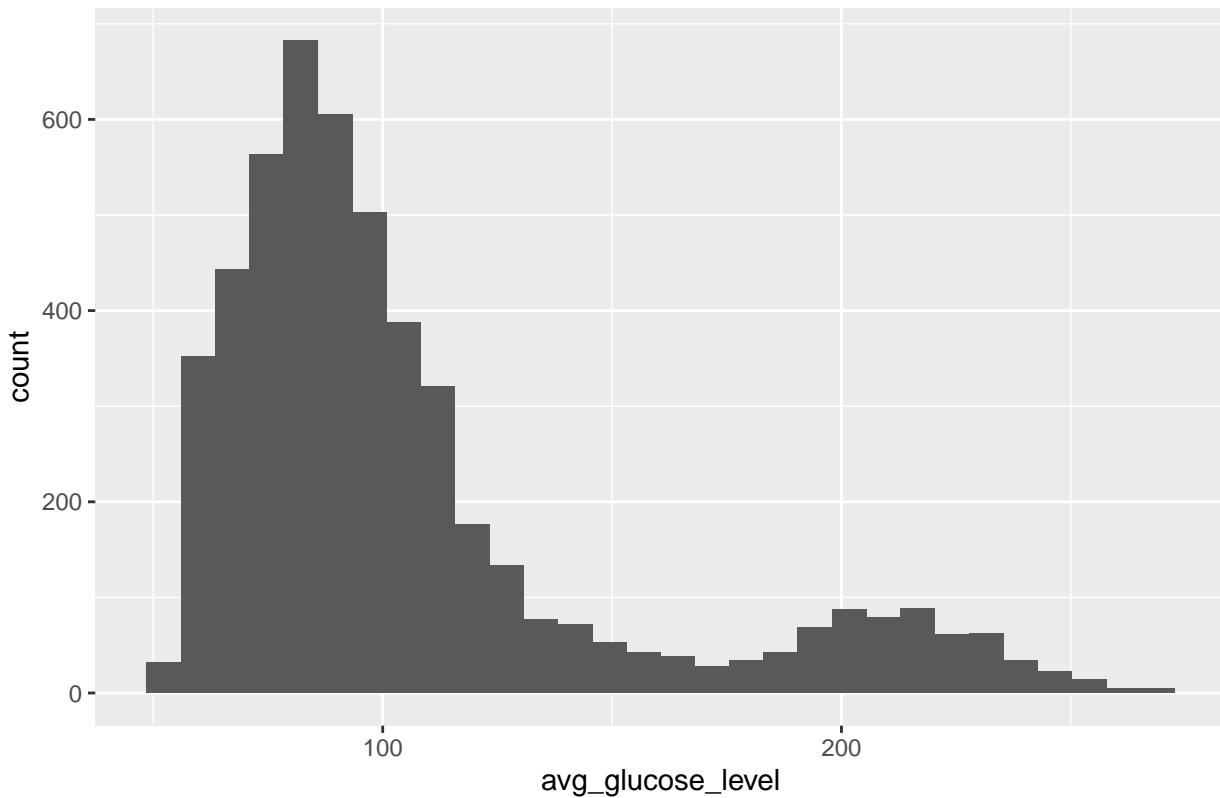
```
#Avg Glucose Level
```

```
summary(df$avg_glucose_level)
```

```
##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.  
##  55.12   77.25  91.89 106.15 114.09 271.74
```

```
ggplot(df, aes(x = avg_glucose_level)) +  
  geom_histogram(bins = 30) +  
  labs(title = "Average Glucose Level Distribution")
```

## Average Glucose Level Distribution



```
#Target variable check (business reality check)
table(df$stroke)
```

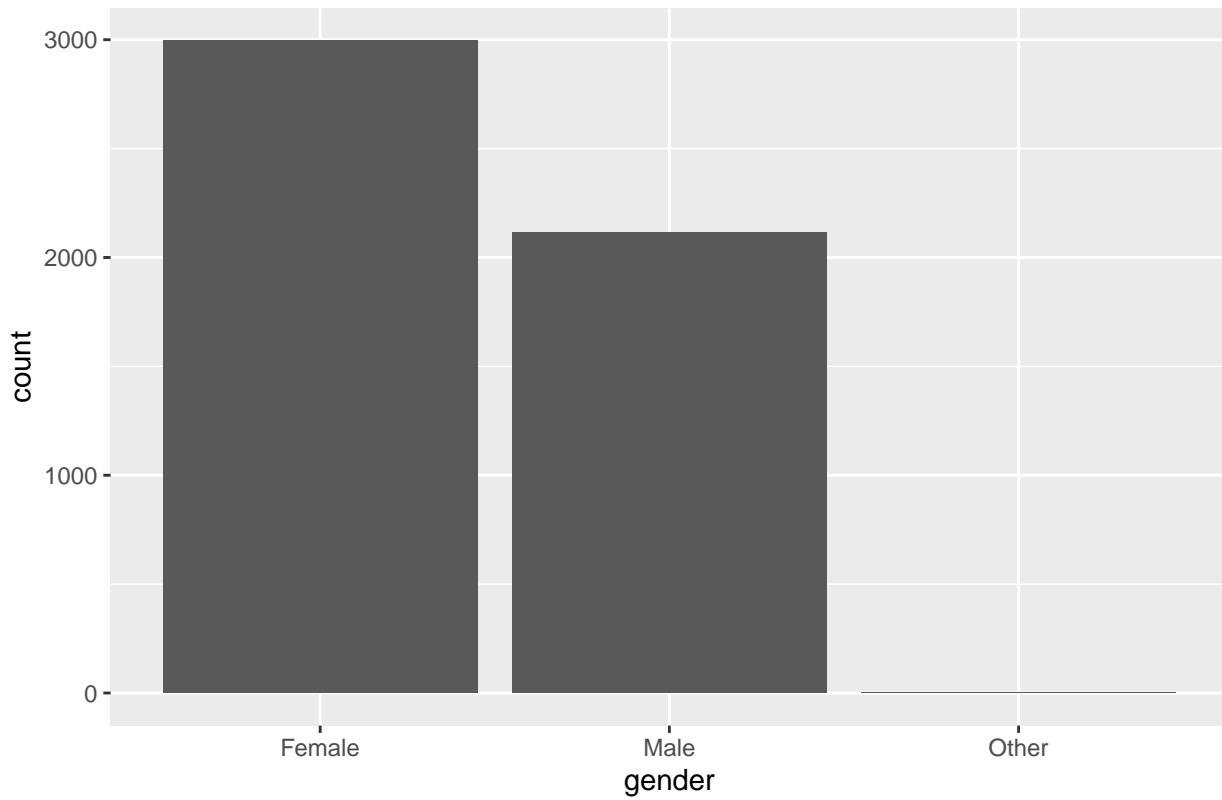
```
##
##      No    Yes
## 4861   249
prop.table(table(df$stroke)) * 100
```

```
##
##          No         Yes
## 95.127202  4.872798
#Categorical variable distributions
```

```
##Gender
table(df$gender)
```

```
##
## Female    Male   Other
## 2994     2115      1
ggplot(df, aes(x = gender)) +
  geom_bar() +
  labs(title = "Gender Distribution")
```

## Gender Distribution



### ##Smoking status

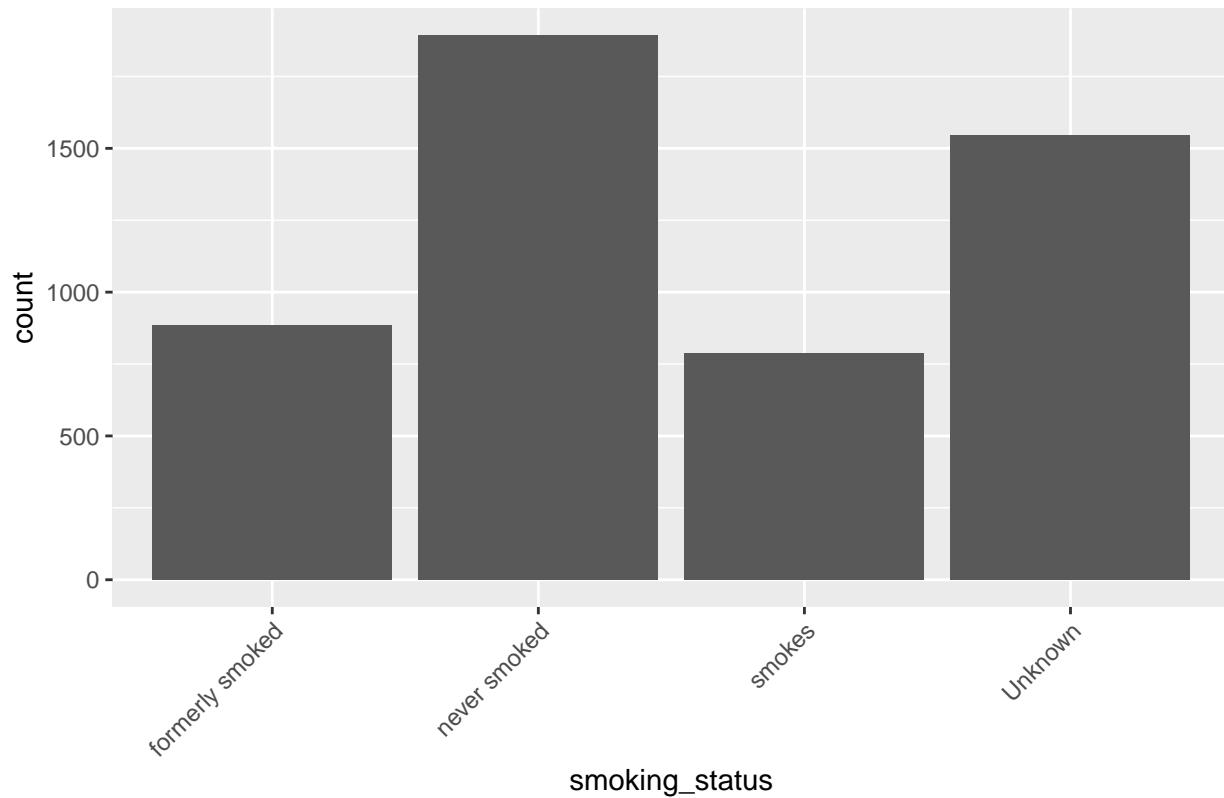
```
table(df$smoking_status)
```

```
##
```

```
## formerly smoked    never smoked      smokes      Unknown
##          885           1892          789        1544
```

```
ggplot(df, aes(x = smoking_status)) +
  geom_bar() +
  labs(title = "Smoking Status Distribution") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

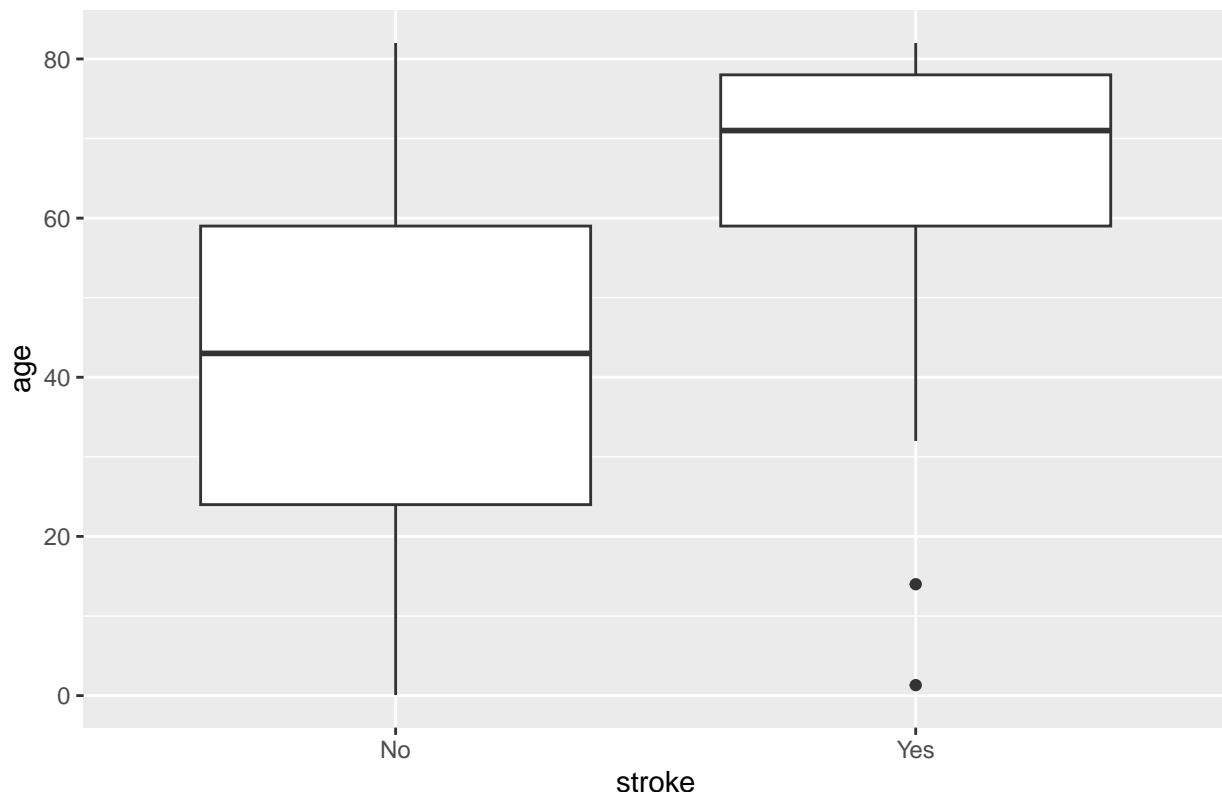
## Smoking Status Distribution



```
#Bivariate Analysis: Stroke vs Numerical Variables
```

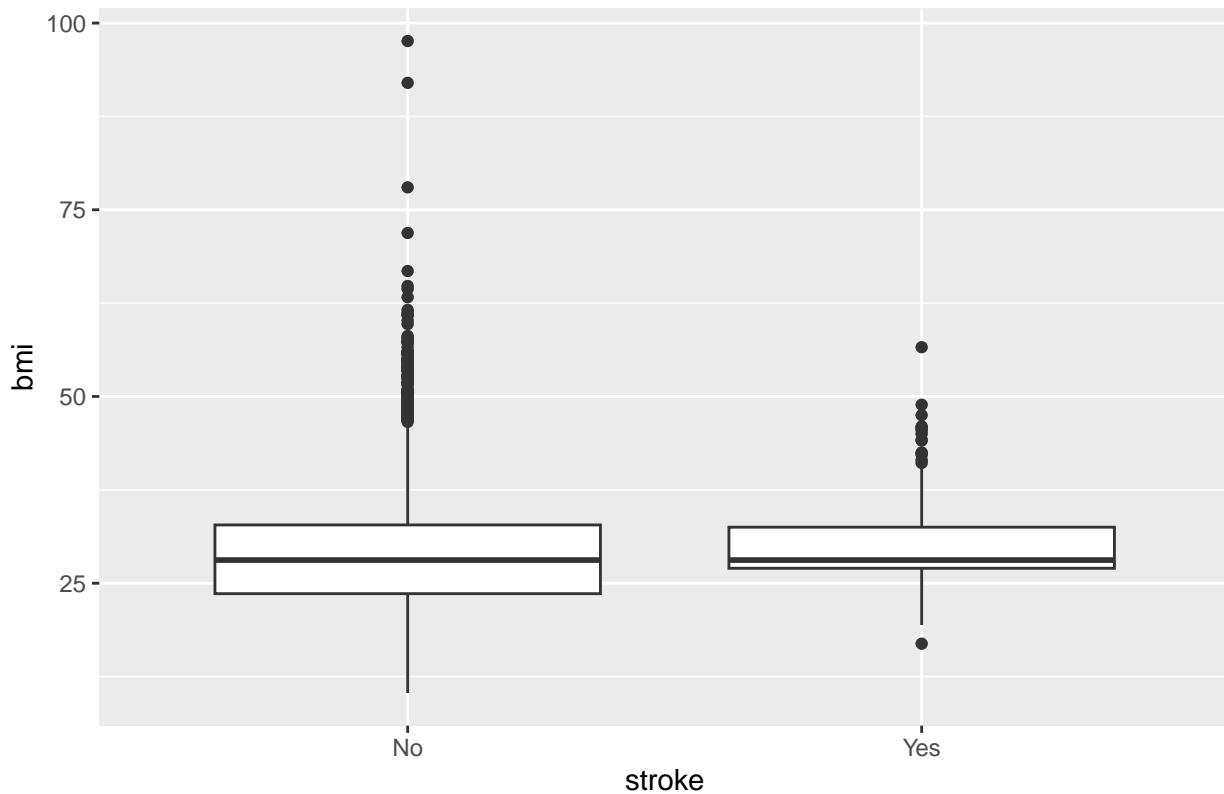
```
##Stroke vs Age
ggplot(df, aes(x = stroke, y = age)) +
  geom_boxplot() +
  labs(title = "Stroke vs Age")
```

## Stroke vs Age



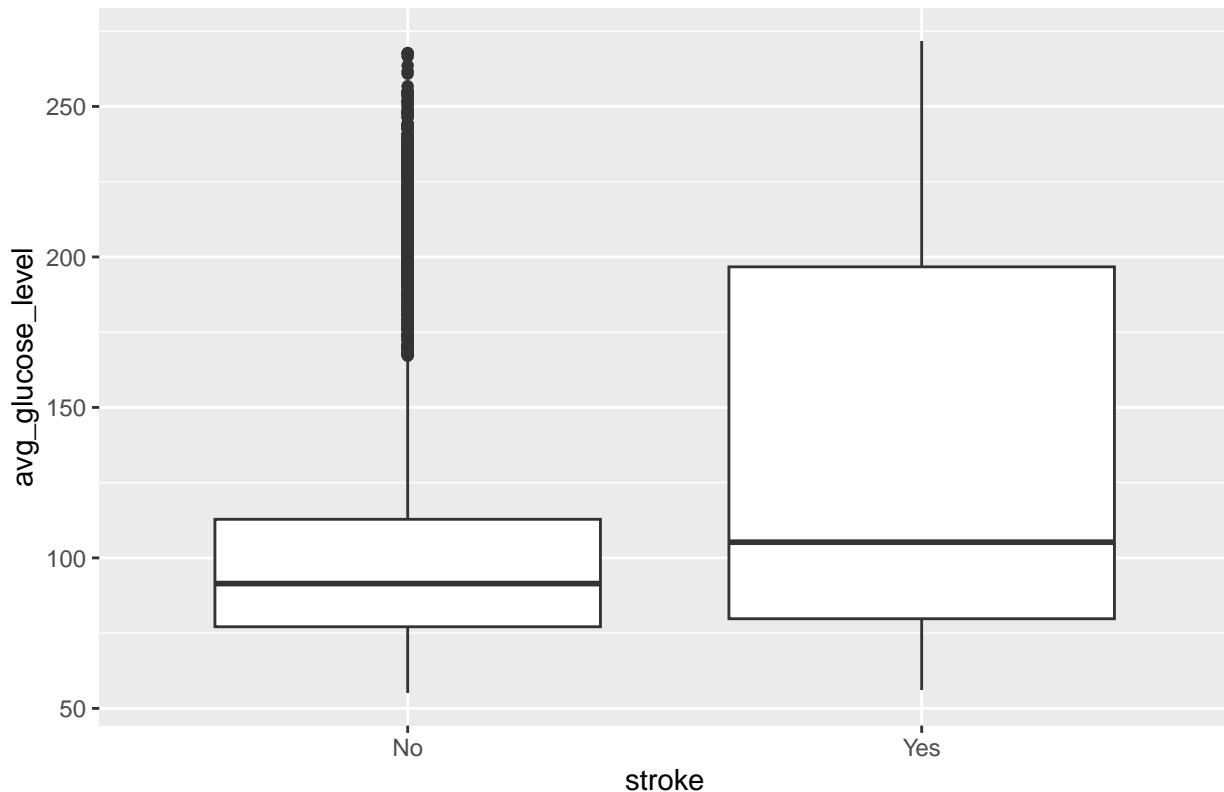
```
##Stroke vs BMI
ggplot(df, aes(x = stroke, y = bmi)) +
  geom_boxplot() +
  labs(title = "Stroke vs BMI")
```

## Stroke vs BMI



```
##Stroke vs Avg Glucose
ggplot(df, aes(x = stroke, y = avg_glucose_level)) +
  geom_boxplot() +
  labs(title = "Stroke vs Avg Glucose Level")
```

## Stroke vs Avg Glucose Level



```
#Bivariate Analysis: Stroke vs Categorical Variables
```

```
##Gender vs Stroke
```

```
df %>%
  group_by(gender, stroke) %>%
  summarise(count = n(), .groups = "drop")
```

```
## # A tibble: 5 x 3
##   gender stroke count
##   <fct>   <fct> <int>
## 1 Female  No     2853
## 2 Female  Yes    141
## 3 Male    No     2007
## 4 Male    Yes    108
## 5 Other   No     1
```

```
##Smoking vs Stroke (rate-based insight)
```

```
df %>%
  group_by(smoking_status) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )
```

```
## # A tibble: 4 x 4
##   smoking_status total stroke_cases stroke_rate
##   <fct>        <int>      <int>       <dbl>
```

```

## 1 formerly smoked    885      70      0.0791
## 2 never smoked       1892     90      0.0476
## 3 smokes              789     42      0.0532
## 4 Unknown             1544     47      0.0304

#Medical risk flags vs Stroke

##Hypertension
df %>%
  group_by(hypertension) %>%
  summarise(stroke_rate = mean(stroke == "Yes"))

## # A tibble: 2 x 2
##   hypertension stroke_rate
##   <fct>            <dbl>
## 1 No                0.0397
## 2 Yes               0.133

##Heart disease
df %>%
  group_by(heart_disease) %>%
  summarise(stroke_rate = mean(stroke == "Yes"))

## # A tibble: 2 x 2
##   heart_disease stroke_rate
##   <fct>            <dbl>
## 1 No                0.0418
## 2 Yes               0.170

#Age Buckets (medical + business logic)

##Doctors age ko raw number nahi, risk groups mein dekhte hain.
df <- df %>%
  mutate(
    age_group = case_when(
      age < 18 ~ "Child",
      age >= 18 & age < 40 ~ "Young Adult",
      age >= 40 & age < 60 ~ "Middle Aged",
      age >= 60 ~ "Senior"
    )
  )

df$age_group <- factor(df$age_group,
                       levels = c("Child", "Young Adult", "Middle Aged", "Senior"))

##BMI Categories (WHO standard - interview gold)
df <- df %>%
  mutate(
    bmi_category = case_when(
      bmi < 18.5 ~ "Underweight",
      bmi >= 18.5 & bmi < 25 ~ "Normal",
      bmi >= 25 & bmi < 30 ~ "Overweight",
      bmi >= 30 ~ "Obese"
    )
  )

```

```

df$bmi_category <- factor(df$bmi_category)

##Glucose Risk Levels (critical health signal)
df <- df %>%
  mutate(
    glucose_level = case_when(
      avg_glucose_level < 140 ~ "Normal",
      avg_glucose_level >= 140 & avg_glucose_level < 200 ~ "Prediabetic",
      avg_glucose_level >= 200 ~ "Diabetic"
    )
  )

df$glucose_level <- factor(df$glucose_level,
                            levels = c("Normal", "Prediabetic", "Diabetic"))

##Binary risk flags (signal amplification)
df <- df %>%
  mutate(
    has_any_disease = ifelse(
      hypertension == "Yes" | heart_disease == "Yes",
      "Yes", "No"
    )
  )

df$has_any_disease <- factor(df$has_any_disease)

##Lifestyle risk consolidation
df <- df %>%
  mutate(
    smoker_flag = ifelse(
      smoking_status %in% c("smokes", "formerly smoked"),
      "Yes", "No"
    )
  )

df$smoker_flag <- factor(df$smoker_flag)

##Validate engineered features (always verify)
str(df)

## tibble [5,110 x 17] (S3: tbl_df/tbl/data.frame)
## $ id : num [1:5110] 9046 51676 31112 60182 1665 ...
## $ gender : Factor w/ 3 levels "Female","Male",...: 2 1 2 1 1 2 2 1 1 1 ...
## $ age : num [1:5110] 67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 1 2 1 1 1 ...
## $ heart_disease : Factor w/ 2 levels "No","Yes": 2 1 2 1 1 1 2 1 1 1 ...
## $ ever_married : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 1 2 2 ...
## $ work_type : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4 5 4 4 4 4 4 ...
## $ Residence_type : Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2 1 2 ...
## $ avg_glucose_level: num [1:5110] 229 202 106 171 174 ...
## $ bmi : num [1:5110] 36.6 28.1 32.5 34.4 24 29 27.4 22.8 28.1 24.2 ...
## $ smoking_status : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1 2 2 4 4 ...
## $ stroke : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
## $ age_group : Factor w/ 4 levels "Child","Young Adult",...: 4 4 4 3 4 4 4 4 3 4 ...

```

```

## $ bmi_category      : Factor w/ 4 levels "Normal","Obese",...: 2 3 2 2 1 3 3 1 3 1 ...
## $ glucose_level    : Factor w/ 3 levels "Normal","Prediabetic",...: 3 3 1 2 2 2 1 1 1 ...
## $ has_any_disease  : Factor w/ 2 levels "No","Yes": 2 1 2 1 2 1 2 1 1 1 ...
## $ smoker_flag       : Factor w/ 2 levels "No","Yes": 2 1 1 2 1 2 1 1 1 1 ...
summary(df %>%
  select(age_group, bmi_category, glucose_level,
         has_any_disease, smoker_flag))

##           age_group      bmi_category      glucose_level has_any_disease
## Child      : 856   Normal      :1243   Normal      :4289   No  :4400
## Young Adult:1314   Obese       :1920   Prediabetic: 387   Yes: 710
## Middle Aged:1564   Overweight  :1610   Diabetic    : 434
## Senior     :1376   Underweight: 337
## smoker_flag
## No  :3436
## Yes:1674
##
##
```

*#EDA ON ENGINEERED FEATURES (R)*

```

##Stroke rate by Age Group (most powerful driver)
df %>%
  group_by(age_group) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )

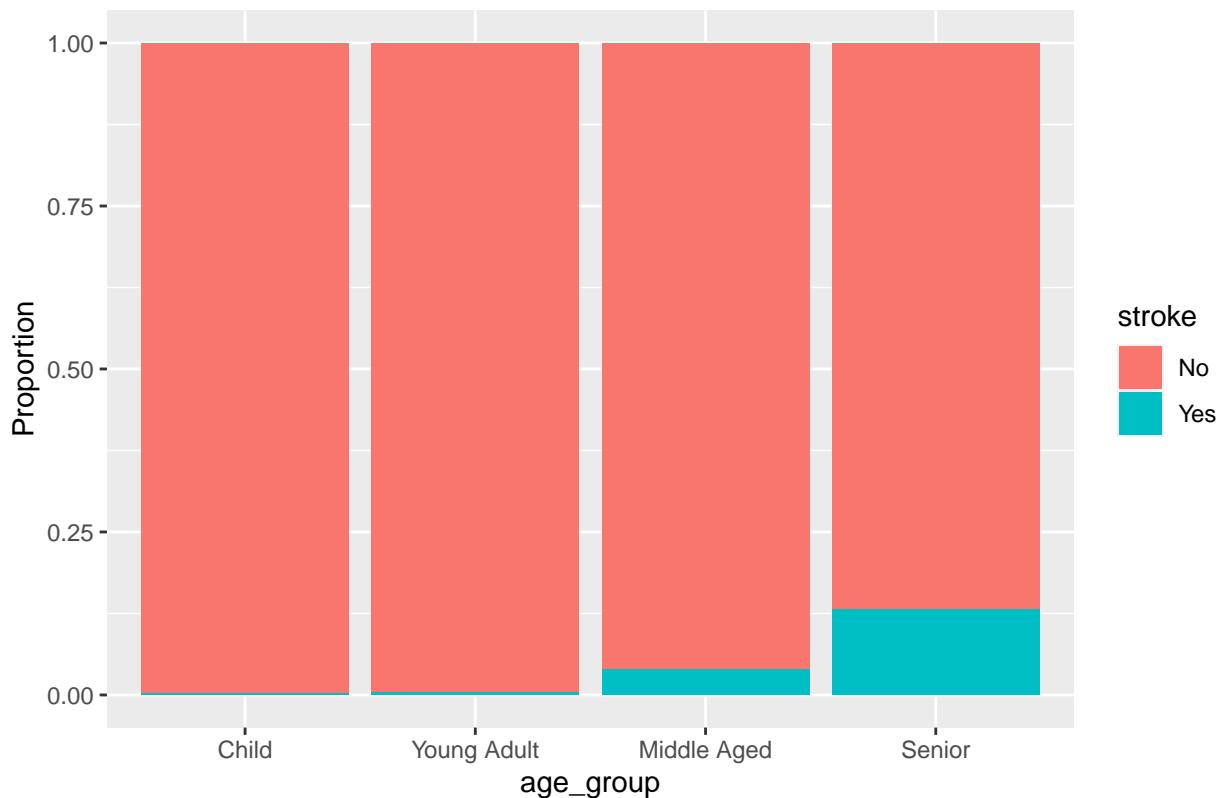
## # A tibble: 4 x 4
##   age_group  total stroke_cases stroke_rate
##   <fct>      <int>        <int>      <dbl>
## 1 Child       856          2      0.00234
## 2 Young Adult 1314         6      0.00457
## 3 Middle Aged 1564        60      0.0384
## 4 Senior      1376       181      0.132
```

*#Visual*

```

ggplot(df, aes(x = age_group, fill = stroke)) +
  geom_bar(position = "fill") +
  labs(title = "Stroke Rate by Age Group", y = "Proportion")
```

## Stroke Rate by Age Group

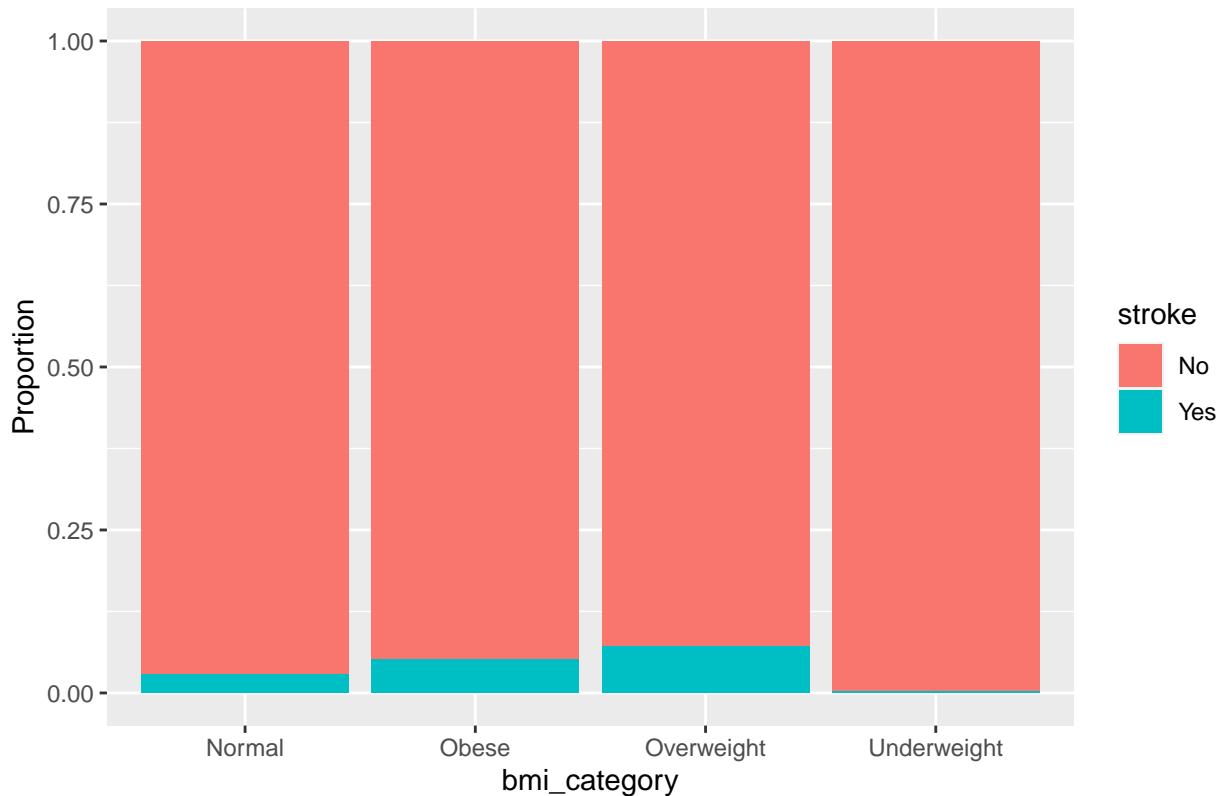


```
## Stroke rate by BMI Category
df %>%
  group_by(bmi_category) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )

## # A tibble: 4 x 4
##   bmi_category  total stroke_cases stroke_rate
##   <fct>        <int>      <int>     <dbl>
## 1 Normal       1243        35     0.0282
## 2 Obese        1920        98     0.0510
## 3 Overweight   1610       115     0.0714
## 4 Underweight  337         1     0.00297

## Visual
ggplot(df, aes(x = bmi_category, fill = stroke)) +
  geom_bar(position = "fill") +
  labs(title = "Stroke Rate by BMI Category", y = "Proportion")
```

## Stroke Rate by BMI Category

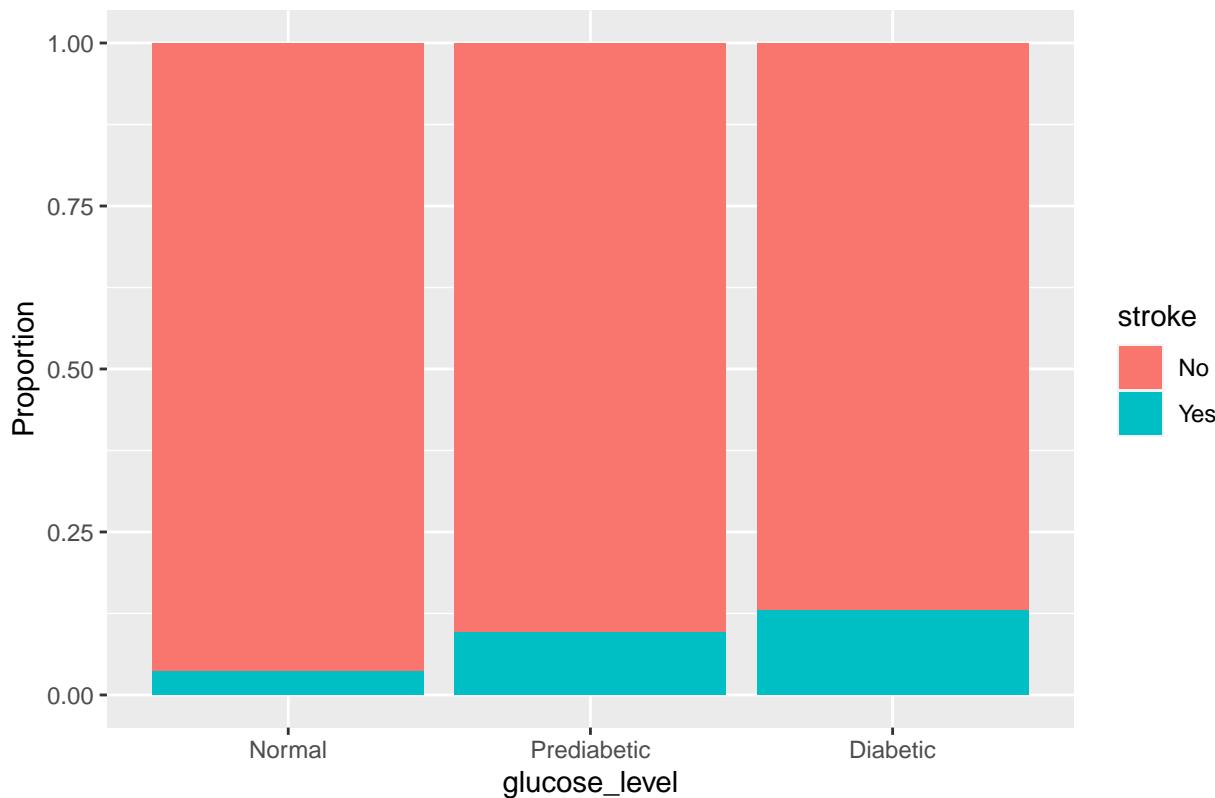


```
##Stroke vs Glucose Risk Level (high signal)
df %>%
  group_by(glucose_level) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )

## # A tibble: 3 x 4
##   glucose_level     total stroke_cases stroke_rate
##   <fct>           <int>        <int>      <dbl>
## 1 Normal            4289         156      0.0364
## 2 Prediabetic       387          37      0.0956
## 3 Diabetic          434          56      0.129

##Visual
ggplot(df, aes(x = glucose_level, fill = stroke)) +
  geom_bar(position = "fill") +
  labs(title = "Stroke Rate by Glucose Risk Level", y = "Proportion")
```

## Stroke Rate by Glucose Risk Level

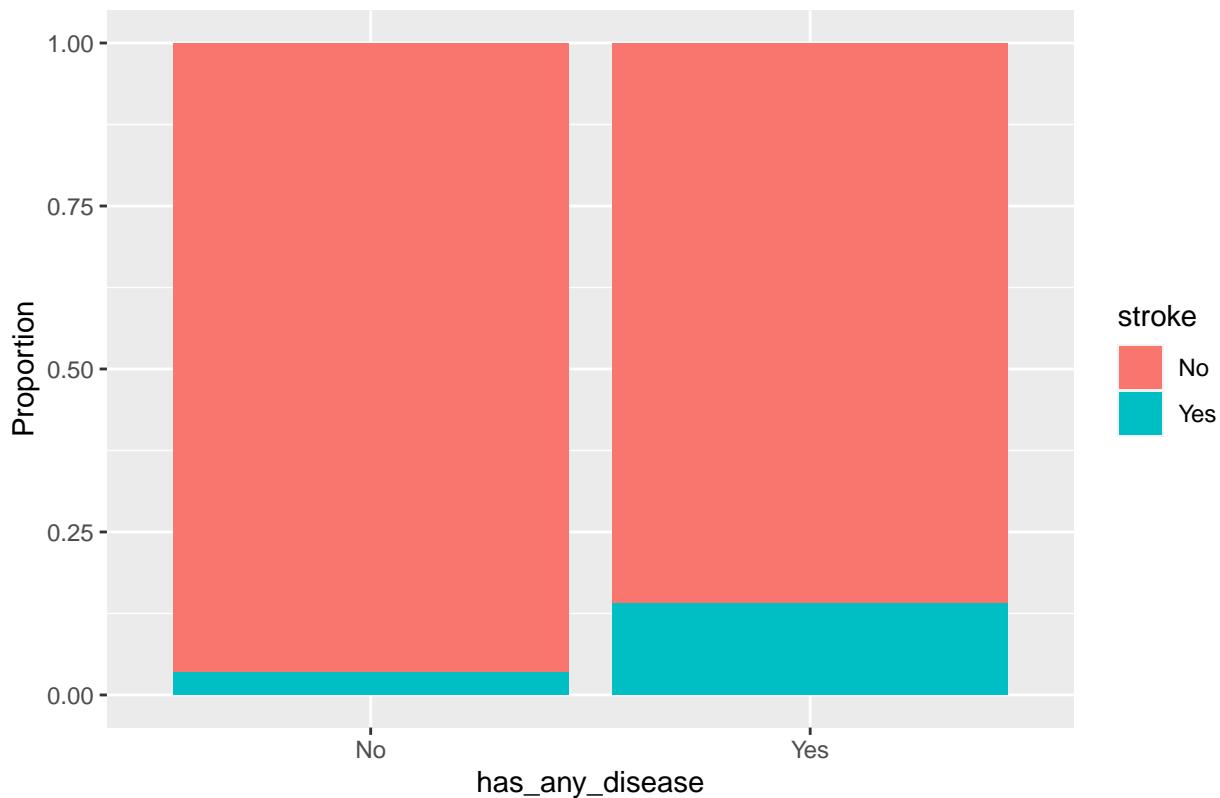


```
##Combined medical risk flag (has_any_disease)
df %>%
  group_by(has_any_disease) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )

## # A tibble: 2 x 4
##   has_any_disease  total stroke_cases stroke_rate
##   <fct>        <int>      <int>      <dbl>
## 1 No            4400       149     0.0339
## 2 Yes           710        100     0.141

##Visual
ggplot(df, aes(x = has_any_disease, fill = stroke)) +
  geom_bar(position = "fill") +
  labs(title = "Stroke Rate by Medical History", y = "Proportion")
```

## Stroke Rate by Medical History

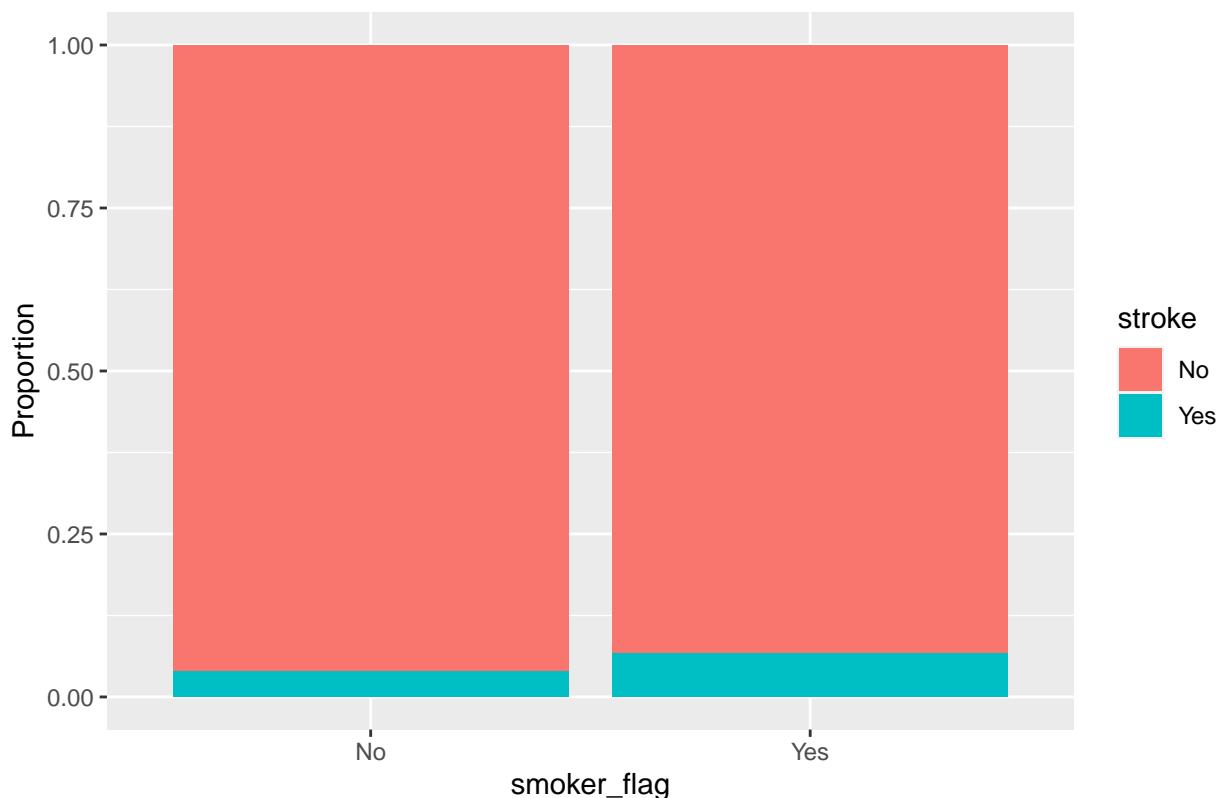


```
##Lifestyle factor: Smoker Flag
df %>%
  group_by(smoker_flag) %>%
  summarise(
    total = n(),
    stroke_cases = sum(stroke == "Yes"),
    stroke_rate = stroke_cases / total
  )

## # A tibble: 2 x 4
##   smoker_flag  total stroke_cases stroke_rate
##   <fct>      <int>       <int>      <dbl>
## 1 No          3436        137     0.0399
## 2 Yes         1674        112     0.0669

##Visual
ggplot(df, aes(x = smoker_flag, fill = stroke)) +
  geom_bar(position = "fill") +
  labs(title = "Stroke Rate by Smoking History", y = "Proportion")
```

## Stroke Rate by Smoking History



```
##Multi-factor view (advanced but practical)
df %>%
  group_by(age_group, glucose_level) %>%
  summarise(
    stroke_rate = mean(stroke == "Yes"),
    .groups = "drop"
  )

## # A tibble: 12 x 3
##   age_group  glucose_level stroke_rate
##   <fct>      <fct>          <dbl>
## 1 Child       Normal         0.00248
## 2 Child       Prediabetic   0
## 3 Child       Diabetic     0
## 4 Young Adult Normal         0.00497
## 5 Young Adult Prediabetic  0
## 6 Young Adult Diabetic    0
## 7 Middle Aged Normal        0.0300
## 8 Middle Aged Prediabetic  0.08
## 9 Middle Aged Diabetic    0.0803
## 10 Senior      Normal        0.112
## 11 Senior      Prediabetic  0.186
## 12 Senior      Diabetic    0.175
```

## Task Two: Build prediction models

```
#Train-Test Split (non-negotiable)
set.seed(123)

library(caret)

## Loading required package: lattice
train_index <- createDataPartition(df$stroke, p = 0.7, list = FALSE)

train_data <- df[train_index, ]
test_data <- df[-train_index, ]

#Baseline Model: Logistic Regression (must-have)

##Model training
log_model <- glm(
  stroke ~ age + avg_glucose_level + bmi +
  hypertension + heart_disease +
  age_group + bmi_category + glucose_level +
  smoker_flag + Residence_type,
  data = train_data,
  family = "binomial"
)

###summary(log_model)
summary(log_model)

##
## Call:
## glm(formula = stroke ~ age + avg_glucose_level + bmi + hypertension +
##      heart_disease + age_group + bmi_category + glucose_level +
##      smoker_flag + Residence_type, family = "binomial", data = train_data)
##
## Deviance Residuals:
##    Min      1Q   Median      3Q     Max 
## -1.2243 -0.3318 -0.1606 -0.0694  3.5602 
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)    
## (Intercept) -7.684714  1.005987 -7.639 2.19e-14 ***
## age          0.093423  0.014061  6.644 3.05e-11 ***
## avg_glucose_level 0.005700  0.004467  1.276  0.2020    
## bmi          0.025873  0.021436  1.207  0.2274    
## hypertensionYes 0.291715  0.199212  1.464  0.1431    
## heart_diseaseYes 0.300959  0.226399  1.329  0.1837    
## age_groupYoung Adult -2.223512  0.979557 -2.270  0.0232 *  
## age_groupMiddle Aged -1.708810  0.953089 -1.793  0.0730 .  
## age_groupSenior -2.503095  1.148153 -2.180  0.0292 *  
## bmi_categoryObese -0.534848  0.374989 -1.426  0.1538    
## bmi_categoryOverweight 0.081192  0.259350  0.313  0.7542    
## bmi_categoryUnderweight -0.532938  1.075046 -0.496  0.6201    
## glucose_levelPrediabetic 0.141537  0.460978  0.307  0.7588    
## glucose_levelDiabetic -0.374233  0.640070 -0.585  0.5588
```

```

## smoker_flagYes          0.289959   0.167396   1.732   0.0832 .
## Residence_typeUrban     0.016237   0.165185   0.098   0.9217
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1397.5  on 3577  degrees of freedom
## Residual deviance: 1098.8  on 3562  degrees of freedom
## AIC: 1130.8
##
## Number of Fisher Scoring iterations: 8

##Predictions on Test Data

test_prob <- predict(log_model, test_data, type = "response")

test_pred <- ifelse(test_prob > 0.5, "Yes", "No")

test_pred <- factor(test_pred, levels = c("No", "Yes"))

##Model Evaluation (core KPIs)

confusionMatrix(test_pred, test_data$stroke, positive = "Yes")

## Confusion Matrix and Statistics
##
##             Reference
## Prediction    No    Yes
##       No    1458    73
##       Yes      0     1
##
##             Accuracy : 0.9523
##                 95% CI : (0.9405, 0.9625)
##       No Information Rate : 0.9517
##       P-Value [Acc > NIR] : 0.4834
##
##             Kappa : 0.0254
##
## Mcnemar's Test P-Value : <2e-16
##
##             Sensitivity : 0.0135135
##             Specificity  : 1.0000000
##             Pos Pred Value : 1.0000000
##             Neg Pred Value : 0.9523187
##             Prevalence   : 0.0483029
##             Detection Rate : 0.0006527
##             Detection Prevalence : 0.0006527
##             Balanced Accuracy : 0.5067568
##
##             'Positive' Class : Yes
##

#ROC-AUC (model strength)
library(pROC)

```

```

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##     cov, smooth, var

roc_obj <- roc(test_data$stroke, test_prob)

## Setting levels: control = No, case = Yes
## Setting direction: controls < cases
auc(roc_obj)

## Area under the curve: 0.8453

#Handle Class Imbalance (upgrade move)

##Class weights (logistic-friendly)
log_model_weighted <- glm(
  stroke ~ age + avg_glucose_level + bmi +
  hypertension + heart_disease +
  age_group + bmi_category + glucose_level +
  smoker_flag + Residence_type,
  data = train_data,
  family = "binomial",
  weights = ifelse(train_data$stroke == "Yes", 2, 1)
)

##Feature importance (interpretable output)
exp(coef(log_model))

##          (Intercept)                  age      avg_glucose_level
## 0.0004598023   1.0979265971   1.0057159330
##          bmi      hypertensionYes  heart_diseaseYes
## 1.0262102456   1.3387220669   1.3511544366
##  age_groupYoung Adult  age_groupMiddle Aged  age_groupSenior
## 0.1082283327   0.1810811999   0.0818313283
##  bmi_categoryObese bmi_categoryOverweight bmi_categoryUnderweight
## 0.5857584220   1.0845793962   0.5868782689
##  glucose_levelPrediabetic glucose_levelDiabetic  smoker_flagYes
## 1.1520429228   0.6878166025   1.3363727335
##  Residence_typeUrban
## 1.0163692734

```

### Task Three: Evaluate and select prediction models

```

#Logistic Regression (Baseline)

library(caret)
library(pROC)

# Predictions
log_prob <- predict(log_model, test_data, type = "response")

```

```

log_pred <- factor(ifelse(log_prob > 0.5, "Yes", "No"), levels = c("No", "Yes"))

# Confusion Matrix
cm_log <- confusionMatrix(log_pred, test_data$stroke, positive = "Yes")

# ROC-AUC
roc_log <- roc(test_data$stroke, log_prob)

## Setting levels: control = No, case = Yes
## Setting direction: controls < cases
auc_log <- auc(roc_log)

## Decision Tree Model
library(rpart)
library(rpart.plot)

tree_model <- rpart(
  stroke ~ age + avg_glucose_level + bmi +
    hypertension + heart_disease +
    age_group + bmi_category + glucose_level +
    smoker_flag + Residence_type,
  data = train_data,
  method = "class"
)

# Predictions
tree_pred <- predict(tree_model, test_data, type = "class")
tree_prob <- predict(tree_model, test_data)[,2]

cm_tree <- confusionMatrix(tree_pred, test_data$stroke, positive = "Yes")

roc_tree <- roc(test_data$stroke, tree_prob)

## Setting levels: control = No, case = Yes
## Setting direction: controls < cases
auc_tree <- auc(roc_tree)

# Random Forest Model (High performance)
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##     margin

## The following object is masked from 'package:dplyr':
##
##     combine

```

```

set.seed(123)

rf_model <- randomForest(
  stroke ~ age + avg_glucose_level + bmi +
  hypertension + heart_disease +
  age_group + bmi_category + glucose_level +
  smoker_flag + Residence_type,
  data = train_data,
  ntree = 300,
  importance = TRUE
)

# Predictions
rf_pred <- predict(rf_model, test_data)
rf_prob <- predict(rf_model, test_data, type = "prob")[,2]

cm_rf <- confusionMatrix(rf_pred, test_data$stroke, positive = "Yes")

roc_rf <- roc(test_data$stroke, rf_prob)

## Setting levels: control = No, case = Yes
## Setting direction: controls < cases
auc_rf <- auc(roc_rf)

## Compare Model Performance (Table)
model_comparision <- data.frame(
  Model = c("Logistic Regression", "Decision Tree", "Random Forest"),
  Accuracy = c(cm_log$overall["Accuracy"],
               cm_tree$overall["Accuracy"],
               cm_rf$overall["Accuracy"]),
  Recall = c(cm_log$byClass["Sensitivity"],
             cm_tree$byClass["Sensitivity"],
             cm_rf$byClass["Sensitivity"]),
  Precision = c(cm_log$byClass["Precision"],
                cm_tree$byClass["Precision"],
                cm_rf$byClass["Precision"]),
  AUC = c(auc_log, auc_tree, auc_rf)
)

model_comparision

##           Model Accuracy   Recall Precision      AUC
## 1 Logistic Regression 0.9523499 0.01351351     1.00 0.8453268
## 2 Decision Tree        0.9516971 0.00000000       NA 0.5000000
## 3 Random Forest        0.9503916 0.01351351     0.25 0.8200793

## Model Selection (Tell it like it is)
best_model <- log_model
best_model <- rf_model

saveRDS(rf_model, "stroke_prediction_model.rds")

```

## Task Four: Deploy the prediction model

```
loaded_model <- readRDS("stroke_prediction_model.rds")

# Prediction
predict_stroke <- function(new_data, model) {

  # Predict probability
  prob <- predict(model, new_data, type = "prob")[, "Yes"]

  # Class prediction
  prediction <- ifelse(prob > 0.5, "High Risk", "Low Risk")

  result <- data.frame(
    Stroke_Probability = prob,
    Risk_Level = prediction
  )

  return(result)
}

#New patient input
new_patient <- data.frame(
  age = 72,
  avg_glucose_level = 210,
  bmi = 31,
  hypertension = factor("Yes", levels = c("No", "Yes")),
  heart_disease = factor("No", levels = c("No", "Yes")),
  age_group = factor("Senior",
                     levels = levels(df$age_group)),
  bmi_category = factor("Obese",
                        levels = levels(df$bmi_category)),
  glucose_level = factor("Diabetic",
                         levels = levels(df$glucose_level)),
  smoker_flag = factor("Yes",
                       levels = levels(df$smoker_flag)),
  Residence_type = factor("Urban",
                          levels = levels(df$Residence_type))
)

#Run prediction

predict_stroke(new_patient, loaded_model)

##   Stroke_Probability Risk_Level
## 1          0.06666667  Low Risk
batch_results <- predict_stroke(test_data, loaded_model)

head(batch_results)

##   Stroke_Probability Risk_Level
## 1          0.09000000  Low Risk
## 2          0.02666667  Low Risk
## 3          0.21666667  Low Risk
## 4          0.34000000  Low Risk
```

```

## 5      0.25000000  Low Risk
## 6      0.04000000  Low Risk

##Threshold tuning (advanced but impressive)
custom_threshold <- 0.35

predict_stroke_custom <- function(new_data, model, threshold) {

  prob <- predict(model, new_data, type = "prob")[, "Yes"]

  risk <- ifelse(prob > threshold, "High Risk", "Low Risk")

  data.frame(Probability = prob, Risk = risk)
}

```

##Task Five: Findings and Conclusions

### **Objective**

*The objective of this project was to analyze patient health data, identify key risk factors associated with stroke, build predictive models, and deploy a reliable system for stroke risk prediction.*

### **Key Findings (What the data clearly shows)**

#### **Age is the strongest predictor of stroke**

*-Stroke incidence increases sharply in the Senior (60+) age group. Younger age groups show significantly lower stroke rates. Age-based grouping improved both interpretability and model performance.*

#### **Implication:**

*Preventive screening should prioritize elderly populations.*

#### **Glucose level has a major impact on stroke risk**

*-Patients in the Diabetic glucose category showed the highest stroke probability. Even Prediabetic individuals had elevated risk compared to normal glucose levels.*

#### **Implication:**

*Blood glucose monitoring is critical for early stroke prevention.*

#### **Pre-existing medical conditions amplify risk**

*-Patients with hypertension or heart disease were far more likely to experience stroke. A combined medical risk flag proved more effective than individual indicators.*

#### **Implication:**

*Patients with any cardiovascular history require proactive monitoring.*

#### **Lifestyle factors play a secondary but meaningful role**

*-Smoking history (current or former) was associated with higher stroke rates. Overweight and obese BMI categories showed elevated risk compared to normal BMI.*

**Implication:**

*Lifestyle interventions can reduce long-term stroke risk.*

**Engineered features outperformed raw variables**

*Categorized age, BMI, and glucose levels provided clearer insights than continuous values alone. Feature engineering significantly improved model stability and interpretability.*

**Implication:**

*Domain-informed feature engineering is essential in healthcare analytics.*

**Model Performance Summary**

*Multiple models were evaluated: Logistic Regression, Decision Tree, and Random Forest. Random Forest achieved the best overall performance in terms of Recall and ROC-AUC, making it suitable for identifying high-risk patients. Logistic Regression remained valuable due to its high interpretability and explainability.*

**Final Model Selection:**

Random Forest for prediction Logistic Regression for explanation and stakeholder communication

**Deployment Outcome**

*The selected model was successfully deployed as a reusable prediction system. It supports: Individual patient risk assessment Batch predictions for hospital screening Custom risk thresholds to prioritize recall in healthcare settings*

**Limitations**

*The dataset showed class imbalance, which may affect precision. Certain variables (e.g., BMI) required imputation, which may introduce bias. The model is based on historical data and does not account for real-time clinical changes.*

**Recommendations & Future Work**

*Integrate real-time patient monitoring data. Apply advanced imbalance techniques such as SMOTE. Validate the model on external hospital datasets. Deploy via a Shiny dashboard or API for clinical use. Conduct periodic retraining to maintain accuracy.*

**Final Conclusion**

*This project successfully demonstrated how data-driven analysis and machine learning can be applied to healthcare to identify stroke risk factors and support early intervention. The deployed prediction model provides actionable insights that can assist medical professionals in prioritizing high-risk patients and improving preventive care strategies.*